

# ROC curve, PR curve, and AUCs

May 28, 2022

Thien

# Data

```
> dt <- iris
> table(dt$Species)

    setosa versicolor  virginica 
        50         50         50 
> dt <- dt[(dt$Species=="versicolor") | (dt$Species=="virginica"),c(1:2,5)]
> table(dt$Species)

    setosa versicolor  virginica 
         0         50         50 
> dt$class <- ifelse(dt$Species == "virginica", 1, 0)
> head(dt)
  Sepal.Length Sepal.Width   Species class
51          7.0         3.2 versicolor     0
52          6.4         3.2 versicolor     0
53          6.9         3.1 versicolor     0
54          5.5         2.3 versicolor     0
55          6.5         2.8 versicolor     0
56          5.7         2.8 versicolor     0
```

```
> library(caret)
> index <- createDataPartition(dt$Species, p=0.80, list=FALSE)
Warning message:
In createDataPartition(dt$Species, p = 0.8, list = FALSE) :
  Some classes have no records ( setosa ) and these will be ignored
> trainDT <- dt[index,]
> table(trainDT$Species)

    setosa versicolor  virginica 
         0         40         40 
> testDT <- dt[-index,]
> table(testDT$Species)

    setosa versicolor  virginica 
         0         10         10
```

# Bulid up logistic regression model

```
> logis <- glm(Species ~ Sepal.Length, trainDT, family="binomial")
> summary(logis)
```

Call:  
glm(formula = Species ~ Sepal.Length, family = "binomial", data = trainDT)

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.87194	-0.88394	-0.06085	1.01817	2.29151

Coefficients:

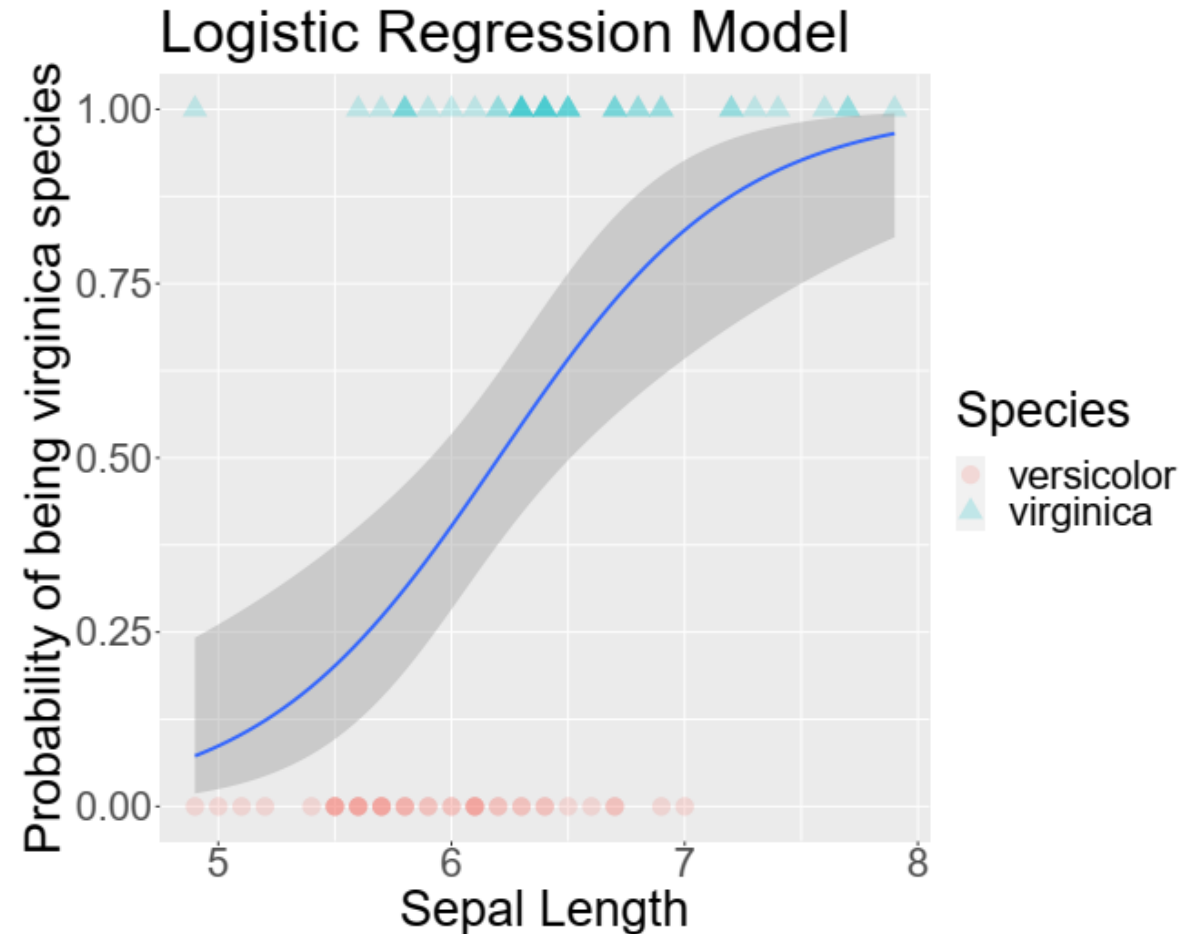
	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-12.1450	3.2479	-3.739	0.000185	***
Sepal.Length	1.9581	0.5237	3.739	0.000185	***

---  
Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 110.904 on 79 degrees of freedom  
Residual deviance: 90.316 on 78 degrees of freedom  
AIC: 94.316

Number of Fisher Scoring iterations: 4



# Prediction

```
> testDT$pred <- predict(logis, testDT, type="response")
> head(testDT)
```


	Sepal.Length	Sepal.Width	Species	class	pred
54	5.5	2.3	versicolor	0	0.2017335
59	6.6	2.9	versicolor	0	0.6853359
61	5.0	2.0	versicolor	0	0.0867056
63	6.0	2.2	versicolor	0	0.4021655
66	6.7	3.1	versicolor	0	0.7259583
70	5.6	2.5	versicolor	0	0.2351086

```
> cutoff <- 0.5
> testDT$pred.species <- ifelse(testDT$pred > cutoff, "virginica",
  "versicolor")
> table(testDT$pred.species)
```

versicolor	virginica
8	12


```
> head(testDT)
```

	Sepal.Length	Sepal.Width	Species	class	pred	pred.species
54	5.5	2.3	versicolor	0	0.2017335	versicolor
59	6.6	2.9	versicolor	0	0.6853359	virginica
61	5.0	2.0	versicolor	0	0.0867056	versicolor
63	6.0	2.2	versicolor	0	0.4021655	versicolor
66	6.7	3.1	versicolor	0	0.7259583	virginica
70	5.6	2.5	versicolor	0	0.2351086	versicolor

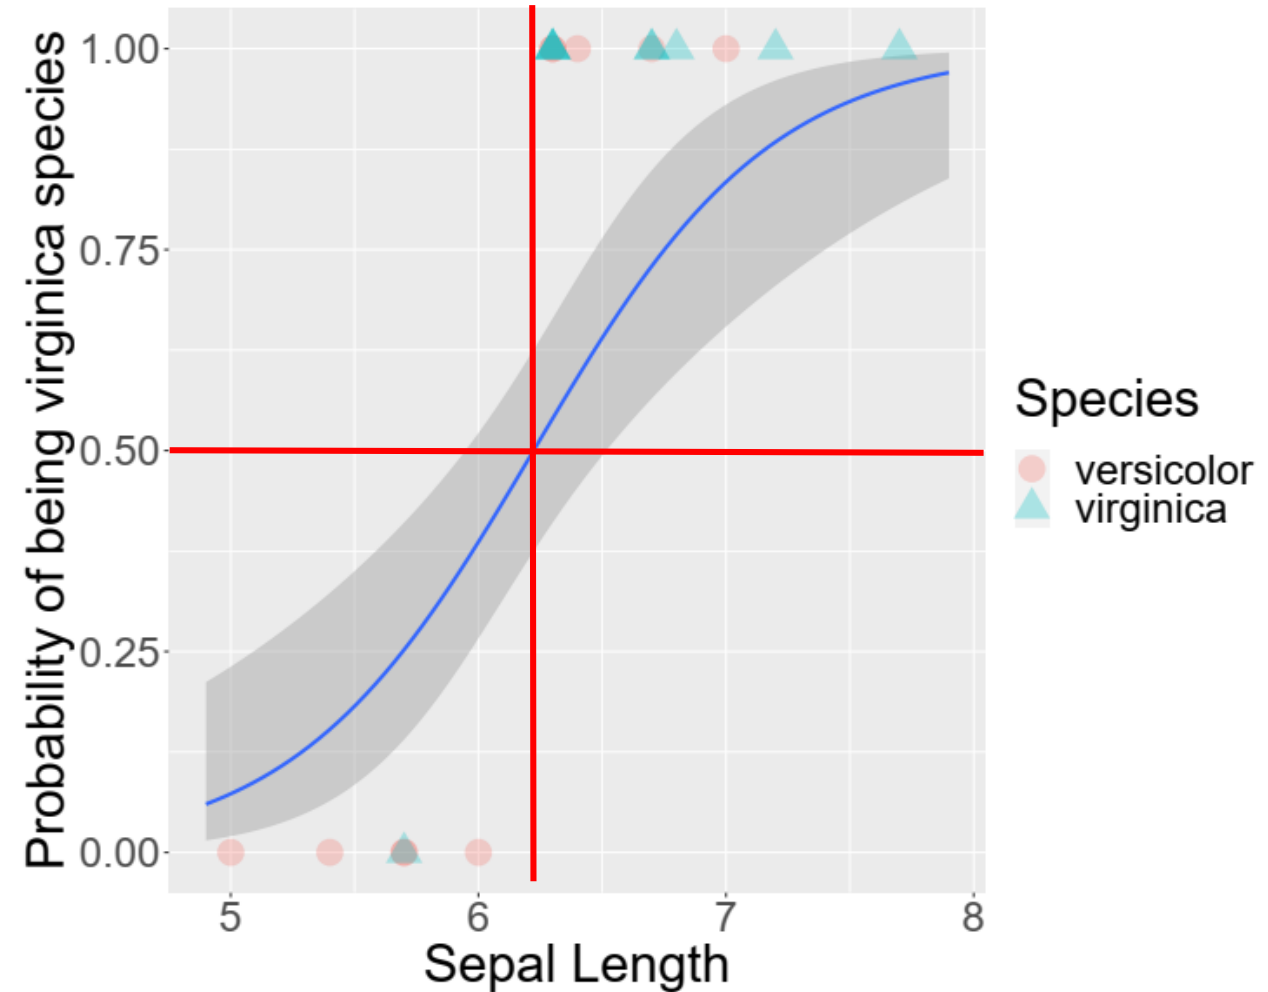


# Threshold = 0.5 for classification

## Original testset



## Predicted testset



# Confusion Matrix

	Virginica	Versicolor
Virginica (Predicted)	9	5
Versicolor (Predicted)	1	5

	Actually Positive (1)	Actually Negative (0)
Predicted Positive (1)	True Positives (TPs)	False Positives (FPs)
Predicted Negative (0)	False Negative (FNs)	True Negative (TNs)

```
> cm <- confusionMatrix(factor(testDT$pred.species),  
+                         factor(testDT$species),  
+                         positive = "virginica")$table  
> cm
```

Prediction	Reference versicolor	virginica
versicolor	5	1
virginica	5	9

# Sensitivity and Specificity

	Actually Positive (1)	Actually Negative (0)
Predicted Positive (1)	True Positives (TPs)	False Positives (FPs)
Predicted Negative (0)	False Negative (FNs)	True Negative (TNs)

- **Sensitivity** (true positive rate): probability of a positive test, conditioned on truly being positive.  $TPR = \frac{TPs}{TPs+FNs}$
- **Specificity** (true negative rate): probability of a negative test, conditioned on truly being negative.  $TNR = \frac{TNs}{FPs+TNs}$

# Sensitivity and Specificity

	Virginica	Versicolor
Virginica (Predicted)	9	5
Versicolor (Predicted)	1	5

- $\text{Sens} = \frac{\text{TPs}}{\text{TPs} + \text{FNs}} = \frac{9}{5 + 9} = 0.9$

- $\text{Spec} = \frac{\text{TNs}}{\text{FPs} + \text{TNs}} = \frac{5}{1 + 5} = 0.5$

```
> cm <- confusionMatrix(factor(testDT$pred.species),  
+                          factor(testDT$species),  
+                          positive = "virginica")$table  
> cm
```


```
          Reference  
Prediction versicolor virginica  
versicolor      5         1  
virginica       5         9
```

```
> # Sens and Spec  
> sensitivity = cm[2, 2] / (cm[2, 2] + cm[1, 2])  
> sensitivity  
[1] 0.9  
> specificity = cm[1, 1] / (cm[1, 1] + cm[2, 1])  
> specificity  
[1] 0.5
```




# Threshold = 1

## Original testset



## Predicted testset



# Threshold = 1

	Virginica	Versicolor
Virginica (Predicted)	0	0
Versicolor (Predicted)	10	10


- $\text{Sens} = \frac{\text{TPs}}{\text{TPs} + \text{FNs}} = \frac{0}{10 + 0} = 0.0$

- $\text{Spec} = \frac{\text{TNs}}{\text{FPs} + \text{TNs}} = \frac{10}{10 + 0} = 1.0$


```
> cm
      Reference
Prediction versicolor virginica
versicolor      10         10
virginica         0          0
> sensitivity = cm[2, 2] / (cm[2, 2] + cm[1, 2])
> sensitivity
[1] 0
> specificity = cm[1, 1] / (cm[1, 1] + cm[2, 1])
> specificity
[1] 1
```

# Threshold = 0

## Original testset



## Predicted testset



# Threshold = 0

	Virginica	Versicolor
Virginica (Predicted)	10	10
Versicolor (Predicted)	0	0

- $\text{Sens} = \frac{\text{TPs}}{\text{TPs} + \text{FNs}} = \frac{10}{0 + 10} = 1.0$

- $\text{Spec} = \frac{\text{TNs}}{\text{FPs} + \text{TNs}} = \frac{0}{0 + 10} = 0.0$

```
> cm
      Reference
Prediction versicolor virginica
versicolor      0         0
virginica      10        10
> sensitivity = cm[2, 2] / (cm[2, 2] + cm[1, 2])
> sensitivity
[1] 1
> specificity = cm[1, 1] / (cm[1, 1] + cm[2, 1])
> specificity
[1] 0
```

**False Positive Rate:** probability of a positive test, conditioned on truly being negative.

Threshold	Sens	Spec
0.0	1.0	0.0
0.1	1.0	0.1
0.2	1.0	0.2
0.3	0.9	0.4
0.4	0.9	0.5
0.5	0.9	0.5
0.6	0.5	0.8
0.7	0.5	0.8
0.8	0.2	0.9
0.9	0.1	1.0
1.0	0.0	1.0

$$\begin{aligned}\text{FPR} &= 1 - \text{TNR} \\ &= \frac{\text{FPs}}{\text{FPs} + \text{TNs}}\end{aligned}$$



FPR
1.0
0.9
0.8
0.6
0.5
0.5
0.2
0.2
0.1
0.0
0.0


# ROC curve

## (Receiver Operating Characteristic curve)

- Showing the performance of a classification model at all classification thresholds
- Provide a simple way to summarize information alternative to tons of confusion matrix
- ROC curve is a compromise of
  - False positive rate (FPR): X-axis
  - True positive rate (TPR): Y-axis


# ROC curve

Threshold	Sens	Spec	FPR
0.0	1.0	0.0	1.0
0.1	1.0	0.1	0.9
0.2	1.0	0.2	0.8
0.3	0.9	0.4	0.6
0.4	0.9	0.5	0.5
0.5	0.9	0.5	0.5
0.6	0.5	0.8	0.2
0.7	0.5	0.8	0.2
0.8	0.2	0.9	0.1
0.9	0.1	1.0	0.0
1.0	0.0	1.0	0.0




# ROC curve

Without having to sort through the confusion matrices, I can tell that this threshold...

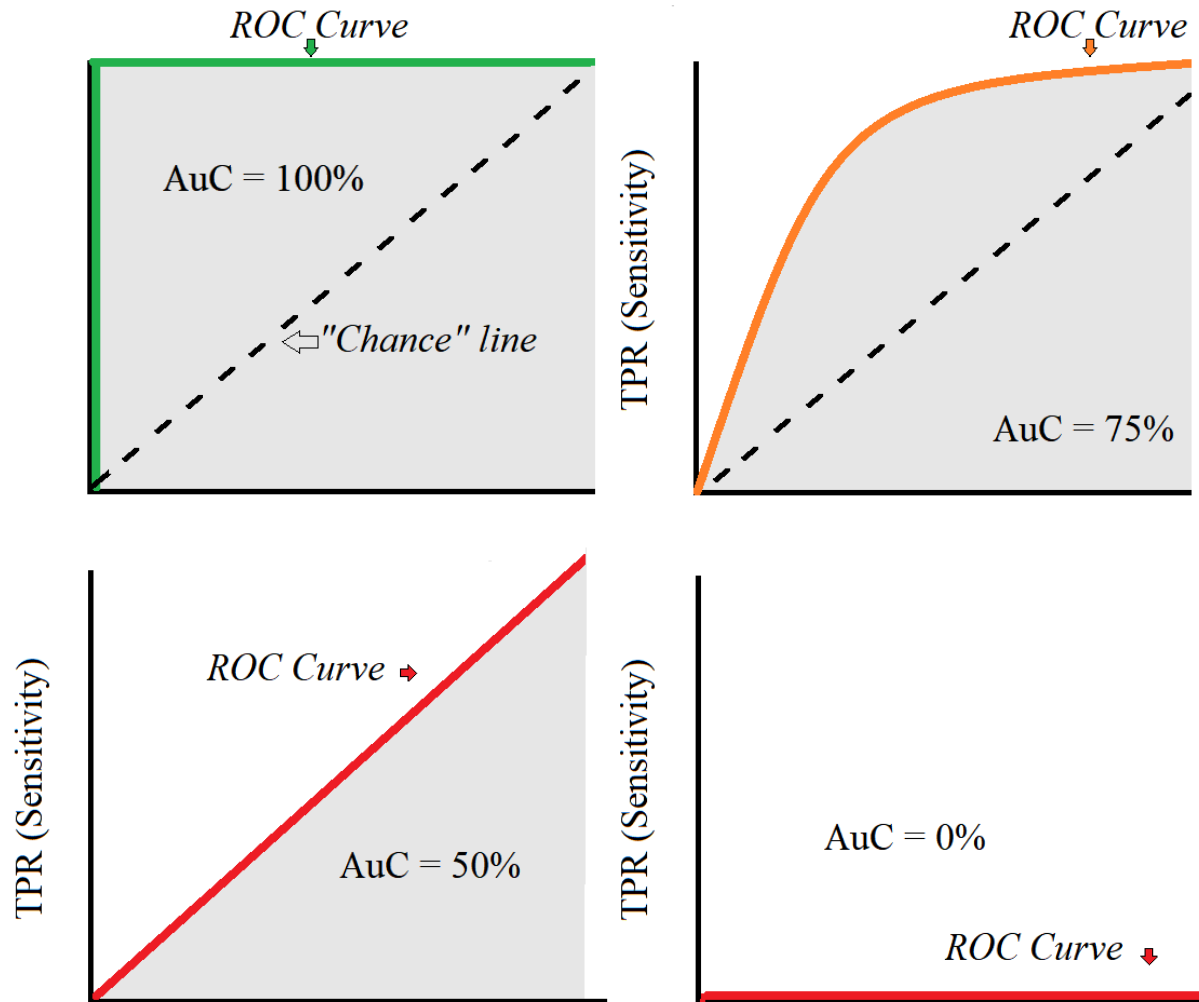


And depending on how many **False Positives** I'm willing to accept, the optimal threshold is either this one...






# AUC (Area Under the Curve)




The **AUC** for the **red** ROC curve is greater than the **AUC** for the **blue** ROC curve, suggesting that the **red** curve is better.



# AUC (Area Under the Curve)

```
> ROC_height = (table_sens_spec$Sens[-1]+table_sens_spec$Sens[-length(table_sens_spec$Sens)])/2  
> ROC_width = -diff(table_sens_spec$FPR)  
> AUC <- sum(ROC_height*ROC_width)  
> AUC  
[1] 0.74
```



**Accuracy:** ratio of number of correct predictions to the total number of observations.

	<b>Actually Positive (1)</b>	<b>Actually Negative (0)</b>
<b>Predicted Positive (1)</b>	True Positives (TPs)	False Positives (FPs)
<b>Predicted Negative (0)</b>	False Negative (FNs)	True Negative (TNs)

- $$\text{Accuracy} = \frac{\text{TPs} + \text{TNs}}{\text{TPs} + \text{TNs} + \text{FPs} + \text{FNs}}$$

# Accuracy

	Virginica	Versicolor
Virginica (Predicted)	9	5
Versicolor (Predicted)	1	5

- Accuracy =  $\frac{TPs+TNs}{TPs+TNs+FPs+FNs} = \frac{9+5}{9+5+5+1} = 0.7$

```
> cm
```

```
Reference
Prediction versicolor virginica
versicolor      5          1
virginica       5          9
```

```
> accuracy = (cm[2, 2] + cm[1, 1]) / (cm[1, 1] + cm[1, 2] + cm[2, 1] + cm[2, 2])
```

```
> accuracy
```

```
[1] 0.7
```

# Accuracy

- What if the test data have 990 Virginica and 10 Versicolor?

	Virginica	Versicolor
Virginica (Predicted)	990	10
Versicolor (Predicted)	0	0

- Accuracy =  $\frac{TPs+TNs}{TPs+TNs+FPs+FNs} = \frac{990+0}{990+0+10+0} = 0.99$
- The accuracy now is 0.99 but completely miss classified all Versicolor species → Imbalance


# Precision and Recall

	Actually Positive (1)	Actually Negative (0)
Predicted Positive (1)	True Positives (TPs)	False Positives (FPs)
Predicted Negative (0)	False Negative (FNs)	True Negative (TNs)


- Precision (positive predictive value): The proportion of positive tests that were correctly classified.  $PPV = \frac{TPs}{TPs+FPs}$
- Recall = Sensitivity:  $TPR = \frac{TPs}{TPs+FNs}$

# Precision and Recall

Original testset



Predicted testset



# Precision and Recall

	Virginica	Versicolor
Virginica (Predicted)	9	5
Versicolor (Predicted)	1	5

- $PPV = \frac{TPs}{TPs+FPs} = \frac{9}{9+5} = 0.64$


- $TPR = Sens = \frac{TPs}{TPs+FNs} = \frac{9}{9+1} = 0.9$

```
> cm
      Reference
Prediction versicolor virginica
versicolor      5         1
virginica       5         9
> ppv = cm[2, 2] / (cm[2, 2] + cm[2, 1])
> ppv
[1] 0.6428571
> # Sens and Spec
> sensitivity = cm[2, 2] / (cm[2, 2] + cm[1, 2])
> sensitivity
[1] 0.9
```




# PR curve

Threshold	Sens	PPV
0.0	1.0	0.5000000
0.1	1.0	0.5263158
0.2	1.0	0.5555556
0.3	0.9	0.6000000
0.4	0.9	0.6428571
0.5	0.9	0.6428571
0.6	0.5	0.7142857
0.7	0.5	0.7142857
0.8	0.2	0.6666667
0.9	0.1	1.0000000



# AUPRC (Area Under the Precision-RecallCurve)

```
> PR_height = (table_sens_ppv$PPV[-1]+table_sens_ppv$PPV[-length(table_sens_ppv$PPV)])/2  
> PR_width = -diff(table_sens_ppv$Sens)  
> AUPRC <- sum(PR_height*PR_width)  
> AUPRC  
[1] 0.7196825
```




# Package for ROC, AUC, PR and AUPRC


```
library(precrc)  
ROCnPR <- evalmod(scores = testDT$pred, labels = testDT$species)  
plot(ROCnPR)  
autoplot(ROCnPR)  
autoplot(ROCnPR, "ROC")  
autoplot(ROCnPR, "PRC")  
aucs <- auc(ROCnPR)  
knitr::kable(aucs)
```

modnames	dsids	curvetypes	aucs
m1	1	ROC	0.7300000
m1	1	PRC	0.7288747


ROC - P: 10, N: 10




Precision-Recall - P: 10, N: 10



ROC - P: 10, N: 10




Precision-Recall - P: 10, N: 10

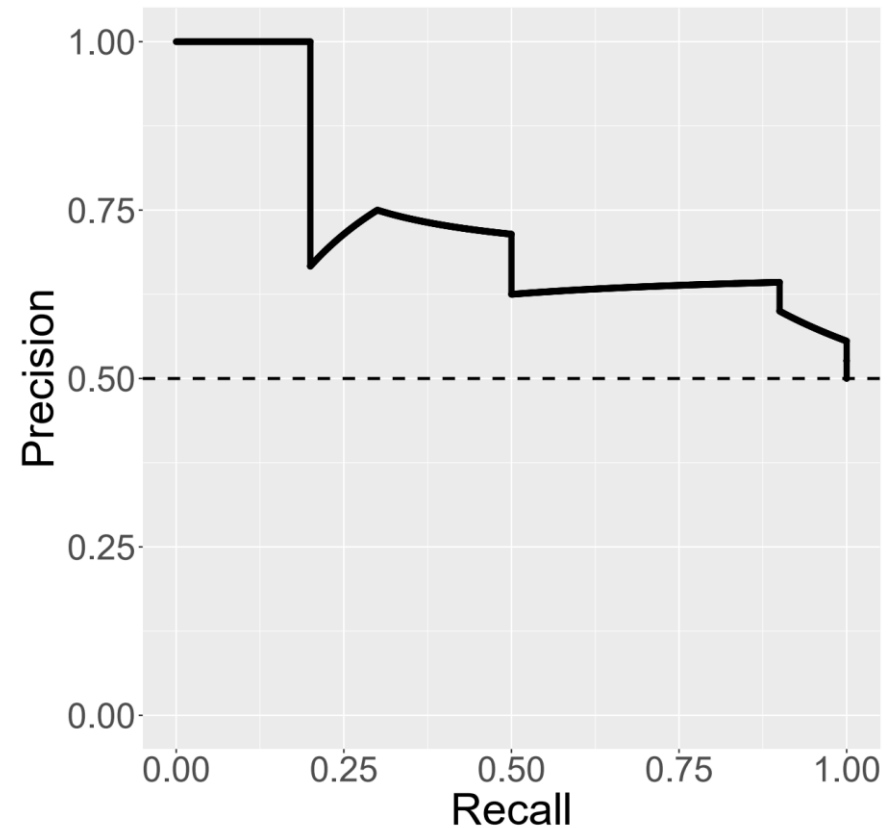


# Package for ROC, AUC, PR and AUPRC

```
dt.roc <- as.data.frame(ROCnPR)[as.data.frame(ROCnPR)$type == "ROC",]  
ggplot(dt.roc, aes(x=x, y=y)) + geom_path(size = 2) + geom_point() +  
  labs(x = "False Positive Rate", y = "True Positive Rate") + xlim(0,1) + ylim(0,1) +  
  geom_abline(intercept= 0, slope = 1, size =1, linetype=2) +  
  theme(text = element_text(size = 30))
```




```
dt.pr <- as.data.frame(ROCnPR)[as.data.frame(ROCnPR)$type == "PRC",]  
ggplot(dt.pr, aes(x=x, y=y)) + geom_path(size = 2) + geom_point() +  
  labs(x = "Recall", y = "Precision") + xlim(0,1) + ylim(0,1) +  
  geom_hline(yintercept = 0.5, size =1, linetype=2) +  
  theme(text = element_text(size = 30))
```



# Decision Boundary

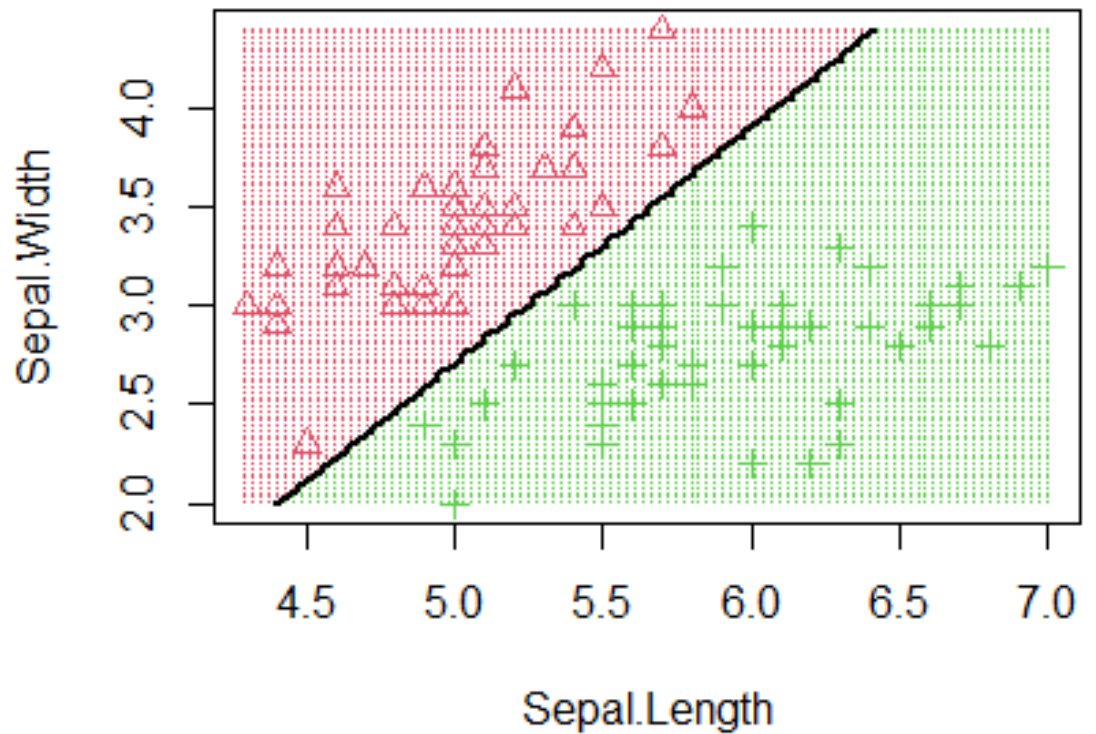
**versicolor / virginica**

**Logistic Regression**



**setosa / versicolor**

**Logistic Regression**



# Exercises

1. Using **Sepal.Length** for **Species (virginica/versicolor)** prediction on test data on multiple thresholds with step=0.01:
  - 1.1. Generate Sensitivity, Specificity, False positive rate and Precision.
  - 1.2. Perform ROC curve, AUC, PR curve and AUPRC.
2. Using **Sepal.Length + Sepal.Width** for **Species (virginica/versicolor)** prediction on test data:
  - 2.1. Perform ROC curve, AUC, PR curve and AUPRC.
  - 2.2. Perform Decision Boundary
3. Using **Sepal.Length + Sepal.Width** for **Species (setosa/versicolor)** prediction on test data on multiple thresholds with step=0.01:
  - 3.1. Perform ROC curve, AUC, PR curve and AUPRC.
  - 3.2. Perform Decision Boundary